

# Homework7\_DS-GA1003

May 6, 2021

Name: Zhuoyuan Xu (Kallen)

NetID: zx1137

## Problem 1-3

```
[ ]: """Computation graph node types

Nodes must implement the following methods:
__init__    - initialize node
forward     - (step 1 of backprop) retrieve output ("out") of predecessor nodes
    ↳ (if
        applicable), update own output ("out"), and set gradient ("d_out")
    ↳ to zero
backward    - (step 2 of backprop), assumes that forward pass has run before.
    Also assumes that backward has been called on all of the node's
    successor nodes, so that self.d_out contains the
    gradient of the graph output with respect to the node output.
    Backward computes summands of the derivative of graph output with
    respect to the inputs of the node, corresponding to paths through
    ↳ the graph
        that go from the node's input through the node to the graph's
    ↳ output.
        These summands are added to the input node's d_out array.
get_predecessors - return a list of the node's parents

Nodes must furthermore have a the following attributes:
node_name    - node's name (a string)
out          - node's output
d_out        - derivative of graph output w.r.t. node output

This computation graph framework was designed and implemented by
Philipp Meerkamp, Pierre Garapon, and David Rosenberg.
License: Creative Commons Attribution 4.0 International License
"""

import numpy as np
```

```

class ValueNode(object):
    """Computation graph node having no input but simply holding a value"""
    def __init__(self, node_name):
        self.node_name = node_name
        self.out = None
        self.d_out = None

    def forward(self):
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        pass

    def get_predecessors(self):
        return []

class VectorScalarAffineNode(object):
    """ Node computing an affine function mapping a vector to a scalar. """
    def __init__(self, x, w, b, node_name):
        """
        Parameters:
        x: node for which x.out is a 1D numpy array
        w: node for which w.out is a 1D numpy array of same size as x.out
        b: node for which b.out is a numpy scalar (i.e. 0dim array)
        node_name: node's name (a string)
        """

        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.x = x
        self.w = w
        self.b = b

    def forward(self):
        self.out = np.dot(self.x.out, self.w.out) + self.b.out
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_x = self.d_out * self.w.out
        d_w = self.d_out * self.x.out
        d_b = self.d_out
        self.x.d_out += d_x
        self.w.d_out += d_w
        self.b.d_out += d_b

```

```

def get_predecessors(self):
    return [self.x, self.w, self.b]

class SquaredL2DistanceNode(object):
    """ Node computing L2 distance (sum of square differences) between 2 arrays.
    → """
    def __init__(self, a, b, node_name):
        """
        Parameters:
        a: node for which a.out is a numpy array
        b: node for which b.out is a numpy array of same shape as a.out
        node_name: node's name (a string)
        """
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.a = a
        self.b = b
        # Variable for caching values between forward and backward
        self.a_minus_b = None

    def forward(self):
        self.a_minus_b = self.a.out - self.b.out
        self.out = np.sum(self.a_minus_b ** 2)
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_a = self.d_out * 2 * self.a_minus_b
        d_b = -self.d_out * 2 * self.a_minus_b
        self.a.d_out += d_a
        self.b.d_out += d_b
        return self.d_out

    def get_predecessors(self):
        return [self.a, self.b]

class L2NormPenaltyNode(object):
    """ Node computing  $l2\_reg * ||w||^2$  for scalars  $l2\_reg$  and vector  $w$  """
    def __init__(self, l2_reg, w, node_name):
        """
        Parameters:
        l2_reg: a numpy scalar array (e.g. np.array(.01)) (not a node)
        w: a node for which w.out is a numpy vector

```

```

        node_name: node's name (a string)
        """
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.l2_reg = np.array(l2_reg)
        self.w = w

    def forward(self):
        self.out = self.l2_reg * np.linalg.norm(self.w.out) ** 2
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_w = self.d_out * (2 * self.l2_reg * self.w.out)
        self.w.d_out += d_w
        return self.d_out

    def get_predecessors(self):
        return [self.w]

class SumNode(object):
    """ Node computing a + b, for numpy arrays a and b """
    def __init__(self, a, b, node_name):
        """
        Parameters:
        a: node for which a.out is a numpy array
        b: node for which b.out is a numpy array of the same shape as a
        node_name: node's name (a string)
        """
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.b = b
        self.a = a

    def forward(self):
        self.out = self.a.out + self.b.out
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_a = self.d_out
        d_b = self.d_out
        self.a.d_out += d_a
        self.b.d_out += d_b

```

```
def get_predecessors(self):
    return [self.a, self.b]
```

```
[ ]: import setup_problem
from sklearn.base import BaseEstimator, RegressorMixin
import numpy as np
import nodes_zx1137 as nodes
import graph
import plot_utils

class RidgeRegression(BaseEstimator, RegressorMixin):
    """ Ridge regression with computation graph """
    def __init__(self, l2_reg=1, step_size=.005, max_num_epochs = 5000):
        self.max_num_epochs = max_num_epochs
        self.step_size = step_size

        # Build computation graph
        self.x = nodes.ValueNode(node_name="x") # to hold a vector input
        self.y = nodes.ValueNode(node_name="y") # to hold a scalar response
        self.w = nodes.ValueNode(node_name="w") # to hold the parameter vector
        self.b = nodes.ValueNode(node_name="b") # to hold the bias parameter
        ↪(scalar)
        self.prediction = nodes.VectorScalarAffineNode(x=self.x, w=self.w,
        ↪b=self.b,
                                                    node_name="prediction")

        # Build computation graph
        self.reg_w = nodes.L2NormPenaltyNode(l2_reg=l2_reg, w=self.w,
        ↪node_name="reg_w")
        self.obj = nodes.SquaredL2DistanceNode(a=self.y, b=self.prediction,
        ↪node_name="diff")
        self.reg_obj = nodes.SumNode(self.reg_w, self.obj, node_name="reg_obj")
        self.graph = graph.ComputationGraphFunction(inputs=[self.x],
        ↪outcomes=[self.y], parameters=[self.w,self.b], prediction=self.prediction,
        ↪objective=self.reg_obj)

    def fit(self, X, y):
        num_instances, num_ftrs = X.shape
        y = y.reshape(-1)

        init_parameter_values = {"w": np.zeros(num_ftrs), "b": np.array(0.0)}
        self.graph.set_parameters(init_parameter_values)

        for epoch in range(self.max_num_epochs):
            shuffle = np.random.permutation(num_instances)
            epoch_obj_tot = 0.0
            for j in shuffle:
```

```

        obj, grads = self.graph.get_gradients(input_values = {"x":X[j]},
        outcome_values = {"y":y[j]})

        #print(obj)
        epoch_obj_tot += obj
        # Take step in negative gradient direction
        steps = {}
        for param_name in grads:
            steps[param_name] = -self.step_size * grads[param_name]
        self.graph.increment_parameters(steps)

    if epoch % 50 == 0:
        train_loss = sum((y - self.predict(X,y)) **2)/num_instances
        print("Epoch ",epoch,": Ave objective=",epoch_obj_tot/
        num_instances," Ave training loss: ",train_loss)

    def predict(self, X, y=None):
        try:
            getattr(self, "graph")
        except AttributeError:
            raise RuntimeError("You must train classifier before predicting data!")

        num_instances = X.shape[0]
        preds = np.zeros(num_instances)
        for j in range(num_instances):
            preds[j] = self.graph.get_prediction(input_values={"x":X[j]})

        return preds

def main():
    data_fname = "data.pickle"
    x_train, y_train, x_val, y_val, target_fn, coefs_true, featurize =
    setup_problem.load_problem(data_fname)

    # Generate features
    X_train = featurize(x_train)
    X_val = featurize(x_val)

    pred_fns = []
    x = np.sort(np.concatenate([np.arange(0,1,.001), x_train]))
    X = featurize(x)

    l2reg = 1

```

```

    estimator = RidgeRegression(l2_reg=l2reg, step_size=0.00005,
    ↪max_num_epochs=2000)
    estimator.fit(X_train, y_train)
    name = "Ridge with L2Reg="+str(l2reg)
    pred_fns.append({"name":name, "preds": estimator.predict(X) })

    l2reg = 0
    estimator = RidgeRegression(l2_reg=l2reg, step_size=0.0005,
    ↪max_num_epochs=500)
    estimator.fit(X_train, y_train)
    name = "Ridge with L2Reg="+str(l2reg)
    pred_fns.append({"name":name, "preds": estimator.predict(X) })

    # Let's plot prediction functions and compare coefficients for several fits
    # and the target function.

    pred_fns.append({"name": "Target Parameter Values (i.e. Bayes Optimal)",
    ↪"coefs": coefs_true, "preds": target_fn(x)})

    plot_utils.plot_prediction_functions(x, pred_fns, x_train, y_train,
    ↪legend_loc="best")

if __name__ == '__main__':
    main()

```

```
[1]: %run ridge_regression.t.py
```

```

DEBUG: (Node l2 norm node) Max rel error for partial deriv w.r.t. w is
6.372644513104131e-08.
.DEBUG: (Node sum node) Max rel error for partial deriv w.r.t. a is
2.8755654572009253e-11.
DEBUG: (Node sum node) Max rel error for partial deriv w.r.t. b is
2.8755654572009253e-11.
.DEBUG: (Parameter w) Max rel error for partial deriv 1.0832498418250635e-09.
DEBUG: (Parameter b) Max rel error for partial deriv 6.458361260311016e-10.
.
-----
Ran 3 tests in 0.006s

OK

```

```
[3]: %run ridge_regression.py
logging.getLogger('matplotlib.font_manager').disabled = True
```

```

Epoch 0 : Ave objective= 1.6094806213101527  Ave training loss:
0.8236515849382391
Epoch 50 : Ave objective= 0.3279977959462063  Ave training loss:

```

0.2417651352044313

Epoch 100 : Ave objective= 0.31649599409323076 Ave training loss:  
0.21184374879140772

Epoch 150 : Ave objective= 0.31452890784244747 Ave training loss:  
0.2037184825639585

Epoch 200 : Ave objective= 0.3121485084212521 Ave training loss:  
0.20034182956314953

Epoch 250 : Ave objective= 0.31280373603957234 Ave training loss:  
0.19877625082307304

Epoch 300 : Ave objective= 0.31248714459392785 Ave training loss:  
0.19818356809561696

Epoch 350 : Ave objective= 0.3115031569492875 Ave training loss:  
0.19765508280926666

Epoch 400 : Ave objective= 0.31067210531127976 Ave training loss:  
0.19910064024346144

Epoch 450 : Ave objective= 0.3110554833622063 Ave training loss:  
0.19774998980617353

Epoch 500 : Ave objective= 0.3102123976915999 Ave training loss:  
0.19767856145932328

Epoch 550 : Ave objective= 0.30953596229096914 Ave training loss:  
0.19895346084883084

Epoch 600 : Ave objective= 0.3098737660877904 Ave training loss:  
0.19751472391091557

Epoch 650 : Ave objective= 0.3093150850808455 Ave training loss:  
0.19817908091334666

Epoch 700 : Ave objective= 0.3083909295851297 Ave training loss:  
0.19769534919964862

Epoch 750 : Ave objective= 0.3086036175478192 Ave training loss:  
0.1977093852431293

Epoch 800 : Ave objective= 0.3083578705816468 Ave training loss:  
0.19806412683777122

Epoch 850 : Ave objective= 0.30770838927223004 Ave training loss:  
0.1980111792987666

Epoch 900 : Ave objective= 0.3078063091050808 Ave training loss:  
0.1980079992986024

Epoch 950 : Ave objective= 0.3082859184800666 Ave training loss:  
0.19838950013312925

Epoch 1000 : Ave objective= 0.30731243399444386 Ave training loss:  
0.19856697674326151

Epoch 1050 : Ave objective= 0.30756780593877253 Ave training loss:  
0.19826994478050178

Epoch 1100 : Ave objective= 0.3072806931251411 Ave training loss:  
0.19830281861019144

Epoch 1150 : Ave objective= 0.3069821273999807 Ave training loss:  
0.19820262441581704

Epoch 1200 : Ave objective= 0.30636961118199 Ave training loss:  
0.19858780093191375

Epoch 1250 : Ave objective= 0.30721375873131973 Ave training loss:



0.19893497958744522

Epoch 1300 : Ave objective= 0.3058722447134926 Ave training loss:  
0.199394484728296

Epoch 1350 : Ave objective= 0.30601599052107437 Ave training loss:  
0.19928444465708506

Epoch 1400 : Ave objective= 0.30613201046323657 Ave training loss:  
0.19860988095978852

Epoch 1450 : Ave objective= 0.3057938230657142 Ave training loss:  
0.19875467075644532

Epoch 1500 : Ave objective= 0.3056803249217344 Ave training loss:  
0.19913444351234513

Epoch 1550 : Ave objective= 0.3049570024219264 Ave training loss:  
0.1999538172637944

Epoch 1600 : Ave objective= 0.30532621315949304 Ave training loss:  
0.20000587406509143

Epoch 1650 : Ave objective= 0.30512372479389344 Ave training loss:  
0.19926491574168348

Epoch 1700 : Ave objective= 0.30512865057760874 Ave training loss:  
0.19966016044785803

Epoch 1750 : Ave objective= 0.30522086208902005 Ave training loss:  
0.19964628147040353

Epoch 1800 : Ave objective= 0.3054030612989358 Ave training loss:  
0.1995722944371397

Epoch 1850 : Ave objective= 0.30529700031436746 Ave training loss:  
0.19954450696300843

Epoch 1900 : Ave objective= 0.30397444733770007 Ave training loss:  
0.19955616786571795

Epoch 1950 : Ave objective= 0.3049864462176651 Ave training loss:  
0.19938840951306924

Epoch 0 : Ave objective= 0.7480293137923937 Ave training loss:  
0.40403905283902763

Epoch 50 : Ave objective= 0.12531880760265604 Ave training loss:  
0.11299507360210885

Epoch 100 : Ave objective= 0.09097861100920873 Ave training loss:  
0.13080076322913367

Epoch 150 : Ave objective= 0.0835764081421503 Ave training loss:  
0.11156742506990712

Epoch 200 : Ave objective= 0.07411701764460152 Ave training loss:  
0.07023988961800091

Epoch 250 : Ave objective= 0.07067633602316048 Ave training loss:  
0.05853692783371531


Epoch 300 : Ave objective= 0.0633720787500745 Ave training loss:  
0.054165150329533514

Epoch 350 : Ave objective= 0.05809762372635873 Ave training loss:  
0.051974487926738956

Epoch 400 : Ave objective= 0.05668885490648498 Ave training loss:  
0.04528040794802371

Epoch 450 : Ave objective= 0.051348149162253394 Ave training loss:

0.04328437440225366



#### Problem 4

Given the definition  $y = Wx + b$ , we can write any entry  $y_i$  in  $y$  as  $y_i = W_{ij}x_j + b_i$ . In the derivative

$$\frac{\partial J}{\partial W_{ij}} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} \frac{\partial y_r}{\partial W_{ij}}$$

we can see the part

$$\frac{\partial y_r}{\partial W_{ij}} = \begin{cases} x_j & r = i \\ 0 & \text{else} \end{cases}$$

Therefore, the part of the derivative inside the summation  $\frac{\partial J}{\partial y_r} \frac{\partial y_r}{\partial W_{ij}}$  has non-zero value only when  $r = i$ , i.e.  $y_r = y_i$ .

$$\frac{\partial J}{\partial W_{ij}} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} \frac{\partial y_r}{\partial W_{ij}} = \frac{\partial J}{\partial y_i} x_j$$

#### Problem 5

Based on our result for problem 4, the vectorized expression for  $\frac{\partial J}{\partial W}$  is

$$\frac{\partial J}{\partial W} = \frac{\partial J}{\partial y} x^T$$

where  $\frac{\partial J}{\partial y}$  and  $x$  are column vectors.

### Problem 6

First, the elementwise expression for the derivative of  $y$  with respect to  $x$  can be written as  $\frac{\partial y_r}{\partial x_i} = W_{ri}$ . Then the expression for the  $i$ -th entry  $\frac{\partial J}{\partial x_i}$  is

$$\frac{\partial J}{\partial x_i} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} \frac{\partial y_r}{\partial x_i} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} W_{ri} = \left( \frac{\partial J}{\partial y} \right)^T W_{:,i}$$

where  $W_{:,i}$  denotes the  $i$ -th column in the matrix  $W$ .

From the above expression, we can show that for the entire vector  $x$ ,

$$\frac{\partial J}{\partial x} = W^T \left( \frac{\partial J}{\partial y} \right)$$

### Problem 7

The  $i$ -th entry in  $b$  can be written as

$$\frac{\partial J}{\partial b_i} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} \frac{\partial y_r}{\partial b_i}$$

In this expression,  $\frac{\partial y_r}{\partial b_i} = \begin{cases} 1 & r = i \\ 0 & \text{else} \end{cases}$ , i.e.  $\frac{\partial y_r}{\partial b_i} = \delta_{ri}$ .

Thus the expression can be written as

$$\frac{\partial J}{\partial b_i} = \sum_{r=1}^m \frac{\partial J}{\partial y_r} \frac{\partial y_r}{\partial b_i} = \frac{\partial J}{\partial y_r} \delta_{ri} = \frac{\partial J}{\partial y_i}$$

### Problem 8

Given the rule  $(A \odot B)_i = A_i B_i$ , for this problem, to show  $\frac{\partial J}{\partial A} = \frac{\partial J}{\partial S} \odot \sigma'(A)$ , we prove the derivative of  $J$  with respect to the  $i$ -th entry of  $A$  satisfy our expression, i.e.  $\frac{\partial J}{\partial A_i} = \left( \frac{\partial J}{\partial S} \odot \sigma'(A) \right)_i = \frac{\partial J}{\partial S_i} \sigma'(A_i)$

Based on similar ideas from the previous problem, we can write

$$\frac{\partial J}{\partial A_i} = \sum_j \frac{\partial J}{\partial S_j} \frac{\partial S_j}{\partial A_i} = \sum_j \frac{\partial J}{\partial S_j} \delta_{ij} \sigma'(A_i) = \frac{\partial J}{\partial S_i} \sigma'(A_i)$$

Therefore,  $\frac{\partial J}{\partial A} = \frac{\partial J}{\partial S} \odot \sigma'(A)$

### Problem 9-11

```
[ ]: class AffineNode(object):
    """Node implementing affine transformation (W,x,b)-->Wx+b, where W is a
    ↪matrix,
    and x and b are vectors
    Parameters:
    W: node for which W.out is a numpy array of shape (m,d)
    x: node for which x.out is a numpy array of shape (d)
```

*b: node for which b.out is a numpy array of shape (m) (i.e. vector of length m)*

```
"""
def __init__(self, W, x, b, node_name):
    self.node_name = node_name
    self.out = None
    self.d_out = None
    self.W = W
    self.x = x
    self.b = b

def forward(self):
    self.out = np.dot(self.W.out, self.x.out) + self.b.out
    self.d_out = np.zeros(self.out.shape)
    return self.out

def backward(self):
    d_W = np.outer(self.d_out, self.x.out)
    d_x = np.dot(self.W.out.T, self.d_out)
    d_b = self.d_out
    self.W.d_out += d_W
    self.x.d_out += d_x
    self.b.d_out += d_b

def get_predecessors(self):
    return [self.W, self.x, self.b]

class TanhNode(object):
    """Node tanh(a), where tanh is applied elementwise to the array a
    Parameters:
    a: node for which a.out is a numpy array
    """
    def __init__(self, a, node_name):
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.a = a

    def forward(self):
        self.out = np.tanh(self.a.out)
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_a = self.d_out * (1 - self.out ** 2)
        self.a.d_out += d_a
```

```
def get_predecessors(self):
    return [self.a]
```

```
[ ]: import matplotlib.pyplot as plt
import setup_problem
from sklearn.base import BaseEstimator, RegressorMixin
import numpy as np
import nodes_zx1137 as nodes
import graph
import plot_utils
import pdb
#pdb.set_trace()

class MLPRegression(BaseEstimator, RegressorMixin):
    """ MLP regression with computation graph """
    def __init__(self, num_hidden_units=10, step_size=.005, init_param_scale=0.
    ↪01, max_num_epochs = 5000):
        self.num_hidden_units = num_hidden_units
        self.init_param_scale = init_param_scale
        self.max_num_epochs = max_num_epochs
        self.step_size = step_size

        # Build computation graph
        # data
        self.x = nodes.ValueNode(node_name="x") # to hold a vector input
        self.y = nodes.ValueNode(node_name="y") # to hold a scalar response

        # parameters
        self.W1 = nodes.ValueNode(node_name="W1")
        self.b1 = nodes.ValueNode(node_name="b1")
        self.w2 = nodes.ValueNode(node_name="w2")
        self.b2 = nodes.ValueNode(node_name="b2")

        # hidden layer
        self.L = nodes.AffineNode(W=self.W1, x=self.x, b=self.b1,
    ↪node_name="affine")
        self.h = nodes.TanhNode(a=self.L, node_name="tanh")

        # prediction
        self.f = nodes.VectorScalarAffineNode(x=self.h, w=self.w2, b=self.b2,
    ↪node_name="prediction")

        # objective
        self.J = nodes.SquaredL2DistanceNode(a=self.y, b=self.f,
    ↪node_name="objective")
```

```

        self.graph = graph.ComputationGraphFunction(inputs=[self.x],
→outcomes=[self.y], parameters=[self.W1,self.b1,self.w2,self.b2],
→prediction=self.f, objective=self.J)

    def fit(self, X, y):
        num_instances, num_ftrs = X.shape
        y = y.reshape(-1)
        s = self.init_param_scale
        init_values = {"W1": s * np.random.standard_normal((self.
→num_hidden_units, num_ftrs)),
                        "b1": s * np.random.standard_normal((self.
→num_hidden_units)),
                        "w2": s * np.random.standard_normal((self.
→num_hidden_units)),
                        "b2": s * np.array(np.random.randn())) }

        self.graph.set_parameters(init_values)

        for epoch in range(self.max_num_epochs):
            shuffle = np.random.permutation(num_instances)
            epoch_obj_tot = 0.0
            for j in shuffle:
                obj, grads = self.graph.get_gradients(input_values = {"x":
→X[j]},
                                                        outcome_values = {"y":
→y[j]})

                #print(obj)
                epoch_obj_tot += obj
                # Take step in negative gradient direction
                steps = {}
                for param_name in grads:
                    steps[param_name] = -self.step_size * grads[param_name]
                self.graph.increment_parameters(steps)
                #pdb.set_trace()

            if epoch % 50 == 0:
                train_loss = sum((y - self.predict(X,y)) **2)/num_instances
                print("Epoch ",epoch,": Ave objective=",epoch_obj_tot/
→num_instances," Ave training loss: ",train_loss)

    def predict(self, X, y=None):
        try:
            getattr(self, "graph")
        except AttributeError:

```

```

        raise RuntimeError("You must train classifier before predicting data!")
    ↪")

    num_instances = X.shape[0]
    preds = np.zeros(num_instances)
    for j in range(num_instances):
        preds[j] = self.graph.get_prediction(input_values={"x":X[j]})

    return preds

def main():
    data_fname = "data.pickle"
    x_train, y_train, x_val, y_val, target_fn, coefs_true, featurize = ↪
    ↪setup_problem.load_problem(data_fname)

    # Generate features
    X_train = featurize(x_train)
    X_val = featurize(x_val)

    # Let's plot prediction functions and compare coefficients for several fits
    # and the target function.
    pred_fns = []
    x = np.sort(np.concatenate([np.arange(0,1,.001), x_train]))

    pred_fns.append({"name": "Target Parameter Values (i.e. Bayes Optimal)", ↪
    ↪"coefs": coefs_true, "preds": target_fn(x)})

    estimator = MLPRegression(num_hidden_units=10, step_size=0.001, ↪
    ↪init_param_scale=.0005, max_num_epochs=5000)
    x_train_as_column_vector = x_train.reshape(x_train.shape[0],1) # fit ↪
    ↪expects a 2-dim array
    x_as_column_vector = x.reshape(x.shape[0],1) # fit expects a 2-dim array
    estimator.fit(x_train_as_column_vector, y_train)
    name = "MLP regression - no features"
    pred_fns.append({"name":name, "preds": estimator.
    ↪predict(x_as_column_vector) })
    #plot_utils.plot_prediction_functions(x, pred_fns, x_train, y_train, ↪
    ↪legend_loc="best")
    if (1==1):
        X = featurize(x)
        estimator = MLPRegression(num_hidden_units=10, step_size=0.0005, ↪
    ↪init_param_scale=.01, max_num_epochs=500)
        estimator.fit(X_train, y_train)
        name = "MLP regression - with features"

```

```

    pred_fns.append({"name":name, "preds": estimator.predict(X) })
    plot_utils.plot_prediction_functions(x, pred_fns, x_train, y_train,
    ↪ legend_loc="best")

if __name__ == '__main__':
    main()

```

```
[1]: %run mlp_regression.t.py
```

```

DEBUG: (Node affine) Max rel error for partial deriv w.r.t. W is
2.6727224425812336e-09.
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. x is
1.2462746693581928e-09.
DEBUG: (Node affine) Max rel error for partial deriv w.r.t. b is
1.6365789325759677e-09.
.DEBUG: (Node tanh) Max rel error for partial deriv w.r.t. a is
6.926211828825294e-09.
.DEBUG: (Parameter W1) Max rel error for partial deriv 7.757046786710294e-06.
DEBUG: (Parameter b1) Max rel error for partial deriv 6.48927428756324e-07.
DEBUG: (Parameter w2) Max rel error for partial deriv 1.7955448355231087e-09.
DEBUG: (Parameter b2) Max rel error for partial deriv 5.873348475102835e-10.
.

```

```
-----
Ran 3 tests in 0.010s
```

OK

```
[4]: %run mlp_regression.py
logging.getLogger('matplotlib.font_manager').disabled = True
```

```

Epoch  0 : Ave objective= 3.137516391575231  Ave training loss:
2.7222678638542765
Epoch  50 : Ave objective= 0.9453888583974991  Ave training loss:
0.9435144534800395
Epoch 100 : Ave objective= 0.9450755362110279  Ave training loss:
0.9432050024618686
Epoch 150 : Ave objective= 0.9419951828013727  Ave training loss:
0.9400330084192745
Epoch 200 : Ave objective= 0.9155739653251257  Ave training loss:
0.9128468688299803
Epoch 250 : Ave objective= 0.8241666838576367  Ave training loss:
0.8200423909145509
Epoch 300 : Ave objective= 0.7736395414396828  Ave training loss:
0.769560110586592
Epoch 350 : Ave objective= 0.7669902623365532  Ave training loss:
0.7626854449456703
Epoch 400 : Ave objective= 0.762091631250786  Ave training loss:
0.7578812526409072

```




Epoch 450 : Ave objective= 0.7559249967010055 Ave training loss:  
 0.7518482998376459  
 Epoch 500 : Ave objective= 0.7483226554000626 Ave training loss:  
 0.7446265725733725  
 Epoch 550 : Ave objective= 0.7405672136626323 Ave training loss:  
 0.7365476152473991  
 Epoch 600 : Ave objective= 0.73165475561857 Ave training loss:  
 0.7279107979330908  
 Epoch 650 : Ave objective= 0.7229064748217219 Ave training loss:  
 0.7191580691895192  
 Epoch 700 : Ave objective= 0.7144902765253807 Ave training loss:  
 0.7106201108495039  
 Epoch 750 : Ave objective= 0.7060936010647947 Ave training loss:  
 0.7027625219572831  
 Epoch 800 : Ave objective= 0.6995418717829005 Ave training loss:  
 0.6955890017147723  
 Epoch 850 : Ave objective= 0.693231113223115 Ave training loss:  
 0.6893465435430121  
 Epoch 900 : Ave objective= 0.687716429437604 Ave training loss:  
 0.6837866046880056  
 Epoch 950 : Ave objective= 0.6827330002724764 Ave training loss:  
 0.6788743596875594  
 Epoch 1000 : Ave objective= 0.6782506072414097 Ave training loss:  
 0.6743233601527687  
 Epoch 1050 : Ave objective= 0.6729696040769835 Ave training loss:  
 0.669928603291718  
 Epoch 1100 : Ave objective= 0.6690892556428207 Ave training loss:  
 0.664921001954736  
 Epoch 1150 : Ave objective= 0.6635115290643725 Ave training loss:  
 0.6592624756457803  
 Epoch 1200 : Ave objective= 0.6569512744488532 Ave training loss:  
 0.6524487807364113  
 Epoch 1250 : Ave objective= 0.6479152382126745 Ave training loss:  
 0.6432642247580777  
 Epoch 1300 : Ave objective= 0.6364255201627868 Ave training loss:  
 0.6319892007535117  
 Epoch 1350 : Ave objective= 0.6222730125416462 Ave training loss:  
 0.6169863529711642  
 Epoch 1400 : Ave objective= 0.6040840591315946 Ave training loss:  
 0.5989056144803573  
 Epoch 1450 : Ave objective= 0.5835506401251223 Ave training loss:  
 0.5770374072605491  
 Epoch 1500 : Ave objective= 0.5578523041365594 Ave training loss:  
 0.5509867882316521  
 Epoch 1550 : Ave objective= 0.5262059791758896 Ave training loss:  
 0.522095832649203  
 Epoch 1600 : Ave objective= 0.4974530675529656 Ave training loss:  
 0.4899083365889365

Epoch 1650 : Ave objective= 0.4647539318248511 Ave training loss:  
 0.4568021973757934  
 Epoch 1700 : Ave objective= 0.43434791332907663 Ave training loss:  
 0.42571798180719483  
 Epoch 1750 : Ave objective= 0.4047494869542201 Ave training loss:  
 0.39992613129756394  
 Epoch 1800 : Ave objective= 0.3825566739968535 Ave training loss:  
 0.37478767282545794  
 Epoch 1850 : Ave objective= 0.3640545575789637 Ave training loss:  
 0.3568310761958429  
 Epoch 1900 : Ave objective= 0.34971594408444157 Ave training loss:  
 0.3434060205440905  
 Epoch 1950 : Ave objective= 0.3406576940241557 Ave training loss:  
 0.33342122675835045  
 Epoch 2000 : Ave objective= 0.3331512527150483 Ave training loss:  
 0.3271066381688972  
 Epoch 2050 : Ave objective= 0.3283402212664162 Ave training loss:  
 0.32171194860011454  
 Epoch 2100 : Ave objective= 0.32215689509935175 Ave training loss:  
 0.3164335626951301  
 Epoch 2150 : Ave objective= 0.31776050802016853 Ave training loss:  
 0.31627187077233465  
 Epoch 2200 : Ave objective= 0.31349884552685525 Ave training loss:  
 0.31172163102758416  
 Epoch 2250 : Ave objective= 0.31242785152280933 Ave training loss:  
 0.30737287215794473  
 Epoch 2300 : Ave objective= 0.31098707949874005 Ave training loss:  
 0.3036595300330174  
 Epoch 2350 : Ave objective= 0.3055593084871608 Ave training loss:  
 0.3017341644300186  
 Epoch 2400 : Ave objective= 0.30488093133016003 Ave training loss:  
 0.2990894863443083  
 Epoch 2450 : Ave objective= 0.29919686013255636 Ave training loss:  
 0.30028179999252136  
 Epoch 2500 : Ave objective= 0.29907330116861125 Ave training loss:  
 0.29551111855372736  
 Epoch 2550 : Ave objective= 0.2988025007675715 Ave training loss:  
 0.2933350394657824  
 Epoch 2600 : Ave objective= 0.29779170932477167 Ave training loss:  
 0.2915620758480712  
 Epoch 2650 : Ave objective= 0.2947261176645541 Ave training loss:  
 0.28917277916904516  
 Epoch 2700 : Ave objective= 0.2933139049062625 Ave training loss:  
 0.2880350609466849  
 Epoch 2750 : Ave objective= 0.2912573903645786 Ave training loss:  
 0.2855759136126734  
 Epoch 2800 : Ave objective= 0.29023533639332366 Ave training loss:  
 0.28391572915960905

Epoch 2850 : Ave objective= 0.2892298353882299 Ave training loss:  
 0.28433435587897327  
 Epoch 2900 : Ave objective= 0.2863975240975124 Ave training loss:  
 0.2808192822927724  
 Epoch 2950 : Ave objective= 0.2852310141917871 Ave training loss:  
 0.2793354166133381  
 Epoch 3000 : Ave objective= 0.2815364303349382 Ave training loss:  
 0.2784851668849465  
 Epoch 3050 : Ave objective= 0.28023828125662503 Ave training loss:  
 0.2779909956954594  
 Epoch 3100 : Ave objective= 0.27891149964867334 Ave training loss:  
 0.2750388916060075  
 Epoch 3150 : Ave objective= 0.2789955184929161 Ave training loss:  
 0.2739721488393313  
 Epoch 3200 : Ave objective= 0.27645933391463706 Ave training loss:  
 0.2734579917537181  
 Epoch 3250 : Ave objective= 0.27640491250548643 Ave training loss:  
 0.2710998317934567  
 Epoch 3300 : Ave objective= 0.2748622152923085 Ave training loss:  
 0.2697849726722456  
 Epoch 3350 : Ave objective= 0.2737176553260719 Ave training loss:  
 0.26857714620284606  
 Epoch 3400 : Ave objective= 0.27146613164193767 Ave training loss:  
 0.2689247392246583  
 Epoch 3450 : Ave objective= 0.2696311284706083 Ave training loss:  
 0.26839403720785077  
 Epoch 3500 : Ave objective= 0.2675941039479308 Ave training loss:  
 0.26558854367062484  
 Epoch 3550 : Ave objective= 0.2687487112596385 Ave training loss:  
 0.26435048762064806  
 Epoch 3600 : Ave objective= 0.26590618095853086 Ave training loss:  
 0.2647450379821079  
 Epoch 3650 : Ave objective= 0.2678510398591245 Ave training loss:  
 0.2628692717222511  
 Epoch 3700 : Ave objective= 0.2645966766934721 Ave training loss:  
 0.2615990257556361  
 Epoch 3750 : Ave objective= 0.26316108430692475 Ave training loss:  
 0.26126092293409486  
 Epoch 3800 : Ave objective= 0.26341760786848717 Ave training loss:  
 0.25905380059104904  
 Epoch 3850 : Ave objective= 0.2631277720515623 Ave training loss:  
 0.2579485570525661  
 Epoch 3900 : Ave objective= 0.2612898621623534 Ave training loss:  
 0.2588143092891282  
 Epoch 3950 : Ave objective= 0.2610454177904173 Ave training loss:  
 0.2561806270382172  
 Epoch 4000 : Ave objective= 0.2582925107711596 Ave training loss:  
 0.25539201056734767

Epoch 4050 : Ave objective= 0.2572431033759671 Ave training loss:  
 0.2558102322075641  
 Epoch 4100 : Ave objective= 0.2586842474336194 Ave training loss:  
 0.2535538997499953  
 Epoch 4150 : Ave objective= 0.2576558184767501 Ave training loss:  
 0.2527640328870781  
 Epoch 4200 : Ave objective= 0.2558839732695167 Ave training loss:  
 0.2523845076946139  
 Epoch 4250 : Ave objective= 0.2557948287001455 Ave training loss:  
 0.2511707437296304  
 Epoch 4300 : Ave objective= 0.25469554160312896 Ave training loss:  
 0.25063278955977186  
 Epoch 4350 : Ave objective= 0.25452490241521153 Ave training loss:  
 0.24945726881354094  
 Epoch 4400 : Ave objective= 0.2535865000252626 Ave training loss:  
 0.24878551256798875  
 Epoch 4450 : Ave objective= 0.25290970356054987 Ave training loss:  
 0.24798879683635422  
 Epoch 4500 : Ave objective= 0.25131362156478754 Ave training loss:  
 0.24850547360947003  
 Epoch 4550 : Ave objective= 0.2503602331422513 Ave training loss:  
 0.24699068361944426  
 Epoch 4600 : Ave objective= 0.24963334041114124 Ave training loss:  
 0.24702900876331393  
 Epoch 4650 : Ave objective= 0.24974842190562996 Ave training loss:  
 0.2451883736863967  
 Epoch 4700 : Ave objective= 0.24936992700031674 Ave training loss:  
 0.24448763831093995  
 Epoch 4750 : Ave objective= 0.24903905248332667 Ave training loss:  
 0.24391478582267084  
 Epoch 4800 : Ave objective= 0.24823475752764243 Ave training loss:  
 0.2432244407535074  
 Epoch 4850 : Ave objective= 0.24693002138529907 Ave training loss:  
 0.2425504303178255  
 Epoch 4900 : Ave objective= 0.24657090884684896 Ave training loss:  
 0.24212550482009576  
 Epoch 4950 : Ave objective= 0.2465022711516545 Ave training loss:  
 0.24135035321254816  
 Epoch 0 : Ave objective= 3.206540758880964 Ave training loss:  
 2.7600393335666866  
 Epoch 50 : Ave objective= 0.15007484988457595 Ave training loss:  
 0.14527634097721806  
 Epoch 100 : Ave objective= 0.11680875705102287 Ave training loss:  
 0.10942355055645007  
 Epoch 150 : Ave objective= 0.10051792450477723 Ave training loss:  
 0.08974775819159062  
 Epoch 200 : Ave objective= 0.08195475487224442 Ave training loss:  
 0.07623225666388646

Epoch 250 : Ave objective= 0.07625044503815824 Ave training loss:  
0.07161626750269053  
Epoch 300 : Ave objective= 0.06319844228066172 Ave training loss:  
0.09582916210107607  
Epoch 350 : Ave objective= 0.05629184872435556 Ave training loss:  
0.052125969830460084  
Epoch 400 : Ave objective= 0.05347634134934582 Ave training loss:  
0.04388093422255334  
Epoch 450 : Ave objective= 0.047113424587467284 Ave training loss:  
0.040245012029976275



```

↳
↳
-----
NameError                                Traceback (most recent call↳
↳last)

```

```

<ipython-input-4-9ff4424cf970> in <module>
    1 get_ipython().run_line_magic('run', 'mlp_regression.py')
----> 2 logging.getLogger('matplotlib.font_manager').disabled = True

```

NameError: name 'logging' is not defined

## Problem 12-14

```
[ ]: class SoftmaxNode(object):
    """ Softmax node
        Parameters:
        z: node for which z.out is a numpy array
    """
    def __init__(self, z, node_name):
        self.node_name = node_name
        self.out = None
        self.d_out = None
        self.z = z

    def forward(self):
        total_prob = np.sum([np.exp(self.z.out[k]) for k in range(len(self.z.
        →out))])
        self.out = np.array([np.exp(self.z.out[k])/total_prob for k in
        →range(len(self.z.out))])
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        derivative = np.empty([len(self.z.out), len(self.out)])
        for i in range(len(self.z.out)):
            for j in range(len(self.out)):
                if i == j:
                    derivative[i][j] = self.out[i] * (1-self.out[i])
                else:
                    derivative[i][j] = -self.out[i] * self.out[j]

        d_z = np.dot(self.d_out, derivative)
        self.z.d_out += d_z

    def get_predecessors(self):
        return [self.z]

class NLLNode(object):
    """ Node computing NLL loss between 2 arrays.
        Parameters:
        y_hat: a node that contains all predictions
        y_true: a node that contains all labels
    """
    def __init__(self, y_hat, y_true, node_name):
        self.node_name = node_name
        self.out = None
```

```

        self.d_out = None
        self.y_hat = y_hat
        self.y_true = y_true

    def forward(self):
        self.out = -np.log(self.y_hat.out[self.y_true.out])
        self.d_out = np.zeros(self.out.shape)
        return self.out

    def backward(self):
        d_ytrue = self.d_out * (-self.y_hat.out[self.y_true.out]/np.exp(-self.
→out))
        d_yhat = np.zeros(self.y_hat.out.shape)
        d_yhat[self.y_true.out] = 1
        d_yhat = self.d_out * (-d_yhat/np.exp(-self.out))
        self.y_hat.d_out += d_yhat
        self.y_true.d_out += d_ytrue

    def get_predecessors(self):
        return [self.y_hat, self.y_true]

```

```

[ ]: import setup_problem
from sklearn.base import BaseEstimator, RegressorMixin
try:
    from sklearn.datasets.samples_generator import make_blobs
except:
    from sklearn.datasets import make_blobs
import numpy as np
import nodes_zx1137 as nodes
import graph

def calculate_nll(y_preds, y):
    """
    Function that calculate the average NLL loss
    :param y_preds: N * C probability array
    :param y: N int array
    :return:
    """
    return np.mean(-np.log(y_preds)[np.arange(len(y)),y])

class MulticlassClassifier(BaseEstimator, RegressorMixin):
    """ Multiclass prediction """
    def __init__(self, num_hidden_units=10, step_size=.005, init_param_scale=0.
→01, max_num_epochs = 1000, num_class=3):
        self.num_hidden_units = num_hidden_units
        self.init_param_scale = init_param_scale

```

```

self.max_num_epochs = max_num_epochs
self.step_size = step_size
self.num_class = num_class

# Build computation graph
# data
self.x = nodes.ValueNode(node_name="x") # to hold a vector input
self.y = nodes.ValueNode(node_name="y") # to hold a scalar response

# parameters
self.W1 = nodes.ValueNode(node_name="W1")
self.b1 = nodes.ValueNode(node_name="b1")
self.W2 = nodes.ValueNode(node_name="W2")
self.b2 = nodes.ValueNode(node_name="b2")

# hidden layer
self.L = nodes.AffineNode(W=self.W1, x=self.x, b=self.b1,
↪node_name="affineW1b1")
self.h = nodes.TanhNode(a=self.L, node_name="tanh")
self.z = nodes.AffineNode(W=self.W2, x=self.h, b=self.b2,
↪node_name="affineb2w2")

# prediction
self.f = nodes.SoftmaxNode(z=self.z, node_name="softmax")

# objective
self.J = nodes.NLLNode(y_hat=self.f, y_true=self.y,
↪node_name='objective')

self.graph = graph.ComputationGraphFunction(inputs=[self.x],
↪outcomes=[self.y], parameters=[self.W1,self.b1,self.W2,self.b2],
↪prediction=self.f, objective=self.J)

def fit(self, X, y):
    num_instances, num_ftrs = X.shape
    y = y.reshape(-1)
    s = self.init_param_scale
    init_values = {"W1": s * np.random.standard_normal((self.
↪num_hidden_units, num_ftrs)),
                  "b1": s * np.random.standard_normal((self.
↪num_hidden_units)),
                  "W2": np.random.standard_normal((self.num_class, self.
↪num_hidden_units)),
                  "b2": np.array(np.random.randn(self.num_class)) }
    self.graph.set_parameters(init_values)

```



```

        for epoch in range(self.max_num_epochs):
            shuffle = np.random.permutation(num_instances)
            epoch_obj_tot = 0.0
            for j in shuffle:
                obj, grads = self.graph.get_gradients(input_values = {"x":  

↪X[j]},  

                                                    outcome_values = {"y":  

↪y[j]})

                #print(obj)
                epoch_obj_tot += obj
                # Take step in negative gradient direction
                steps = {}
                for param_name in grads:
                    steps[param_name] = -self.step_size * grads[param_name]
                self.graph.increment_parameters(steps)
                #pdb.set_trace()

            if epoch % 50 == 0:
                train_loss = calculate_nll(self.predict(X,y), y)
                print("Epoch ",epoch," Ave training loss: ",train_loss)

    def predict(self, X, y=None):
        try:
            getattr(self, "graph")
        except AttributeError:
            raise RuntimeError("You must train classifier before predicting data!")
        ↪")

        num_instances = X.shape[0]
        preds = []
        for j in range(num_instances):
            preds.append(self.graph.get_prediction(input_values={"x":X[j]}).  

↪reshape(1,-1))

        return np.concatenate(preds, axis=0)

def main():
    # load the data from HW5
    np.random.seed(2)
    X, y = make_blobs(n_samples=500, cluster_std=.25, centers=np.array([(-3,  

↪1), (0, 2), (3, 1)]))
    training_X = X[:300]
    training_y = y[:300]
    test_X = X[300:]
    test_y = y[300:]

```

```

# train the model
estimator = MulticlassClassifier()
estimator.fit(training_X, training_y)

# report test accuracy
test_acc = np.sum(np.argmax(estimator.predict(test_X), axis=1)==test_y)/
→len(test_y)
print("Test set accuracy = {:.3f}".format(test_acc))

if __name__ == '__main__':
    main()

```

```
[1]: %run multiclass.t.py
```

C:\Users\kalle\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:143:  
FutureWarning: The sklearn.datasets.samples\_generator module is deprecated in  
version 0.22 and will be removed in version 0.24. The corresponding classes /  
functions should instead be imported from sklearn.datasets. Anything that cannot  
be imported from sklearn.datasets is now part of the private API.

```
warnings.warn(message, FutureWarning)
DEBUG: (Node softmax) Max rel error for partial deriv w.r.t. z is
1.5385742524215873e-09.
.DEBUG: (Parameter W1) Max rel error for partial deriv 2.1791491311701626e-07.
DEBUG: (Parameter b1) Max rel error for partial deriv 5.4507464417531275e-08.
DEBUG: (Parameter W2) Max rel error for partial deriv 6.754800815066246e-08.
DEBUG: (Parameter b2) Max rel error for partial deriv 2.205053947010376e-08.
.
```

-----  
Ran 2 tests in 0.012s

OK

```
[2]: %run multiclass.py
```

```

Epoch 0   Ave training loss:  0.10767753468425852
Epoch 50  Ave training loss:  0.0037402729498018897
Epoch 100 Ave training loss:  0.0019509875089186069
Epoch 150 Ave training loss:  0.00131892201003299
Epoch 200 Ave training loss:  0.0009947600104512856
Epoch 250 Ave training loss:  0.0007975221227263997
Epoch 300 Ave training loss:  0.0006649220947379011
Epoch 350 Ave training loss:  0.0005697138957458585
Epoch 400 Ave training loss:  0.0004980771960410216
Epoch 450 Ave training loss:  0.00044225221211177576
Epoch 500 Ave training loss:  0.0003975450315101259
Epoch 550 Ave training loss:  0.0003609495175393885

```

Epoch	600	Ave training loss:	0.0003304520224436157
Epoch	650	Ave training loss:	0.0003046529432352649
Epoch	700	Ave training loss:	0.0002825495526238341
Epoch	750	Ave training loss:	0.0002634047943162139
Epoch	800	Ave training loss:	0.00024666486030361454
Epoch	850	Ave training loss:	0.0002319056839595022
Epoch	900	Ave training loss:	0.0002187970217752761
Epoch	950	Ave training loss:	0.00020707801611844357

Test set accuracy = 1.000